

WRITEUP FOR ASSIGNMENT 2

By Arman Rajesh Ganjoo 2021018

Q1.1

I experimented with Linux thread scheduling in this question by creating 3 threads and giving them different scheduling policies (SCHED_OTHER, SCHED_RR, SCHED_FIFO) along with different priority levels (except for threadA using SCHED_OTHER) for each iteration (total 10 iterations for finding the variance in data). Each thread used the same logical function for computation. All of them counted from 1 to 2^{32} asynchronously with each other.

For timing the threads, I used:

`struct timespec x` -> Found in `<time.h>` header


For modifying the scheduling policy and priority levels of the threads, I used:

(There are found in `<sched.h>` header)

```
pthread_attr_t attrThread;
struct sched_param tsp;
pthread_attr_init(&attrThread);
pthread_attr_setinheritsched(&attrThread, PTHREAD_INHERIT_SCHED);
pthread_attr_setschedpolicy(&attrThread, SCHED_FIFO);
pthread_attr_setschedparam(&attrThread, &tsp);
```

For each iteration, the policy of the threads remaining constant i.e ThreadA was SCHED_OTHER, ThreadB was SCHED_RR, ThreadC was SCHED_FIFO. But the priority level of ThreadB and ThreadC kept on incrementing by a certain factor for which I used the formula -> $SCHED_RR/FIFO_minPriority + iterationNumber * priorityBreak$.

After the whole experimentation, it was clear that threads running with the scheduling policy of SCHED_RR and SCHED_FIFO were much faster than threads running with SCHED_OTHER. The fact that SCHED_RR and SCHED_FIFO policy threads could have higher priority levels also meant that the run-time of these particular threads could be decreased further with an increase in their priority levels.




Q1.2

In this question, I had to create 3 different processes using the `fork()` command found in the `<unistd.h>` header. These processes had different scheduling policies (found in the `<sched.h>` header) in the order -> `proc1: SCHED_OTHER`, `proc2: SCHED_RR`, `proc3: SCHED_FIFO`.

Each of these processes was supposed to run a bash script executable (which is created after running `chmod +x ./scriptname` on the bash script). The bash scripts contain the same content but are executed by their respective procs. The bash scripts were executed using `execl()` command. When these bash scripts were executed, they attempted to compile our custom kernel (which takes a long time!). All these 3 processes ran asynchronously and compiled the custom kernel.

The observations I observed were more or less the same as the ones I noted in Q1.1 where threads were used instead of processes.



Q2

In this particular question, I had to generate a patchfile which could be used to patch a kernel and add the custom syscall I made (kernel_2d_memcpy) to it. Before patching the file though, we had to add the custom made system file to the kernel. This syscall was made to copy a 2D matrix from one memory address to another - in gist, just clone the memory. But this cloning was supposed to be done using the `__copy_from_user()` and `__copy_to_user()` kernel functions. I had to include the `<linux/kernel.h>` header file to create this particular syscall.

After creating the custom syscall, I had to create a Makefile which compiled the file which contained the custom syscall.

After that, I had to add the custom syscall as a new entry to `syscall_64.tbl` (as my machine runs on x86_64 architecture). This `.tbl` file was in the kernel directory.

After that, I just compiled the kernel and it successfully added the custom and new syscall to my kernel.

To create the patchfile, I had to use ``diff -ruN`` command on 3 files:

1. The `syscall_64.tbl`
2. The `kernel_2d_memcpy.c`

3. The Makefile in the kernel directory

This created the patchfile which can be patched onto a kernel using the command
``patch -p0 < patchfile``

Reference: <https://brennan.io/2016/11/14/kernel-dev-ep3/>